

Leaf Disease Detection using Deep Learning

Python 3.8+ PyTorch 1.9+ License MIT Dataset Kaggle

- ⌚ A complete end-to-end deep learning solution for automated plant disease detection

Leveraging simplified AlexNet with early exit mechanisms for efficient and accurate classification

[Features](#) • [Installation](#) • [Usage](#) • [Results](#) • [Documentation](#)

📋 Table of Contents

- [Overview](#)
- [Dataset](#)
- [Features](#)
- [Installation](#)
- [Usage](#)
- [Project Structure](#)
- [Model Architecture](#)
- [Results](#)
- [Requirements](#)

🔍 Overview

This project implements an **automated leaf disease detection system** using state-of-the-art convolutional neural networks (CNN). The system can accurately classify leaf images into **5 disease categories** with high precision.

📊 Disease Classes

 Bacterial Leaf Spot  Downy Mildew  Healthy Leaf

 Mosaic Disease  Powdery Mildew

⌚ Key Capabilities

-  **High Accuracy:** 85-95% validation accuracy
-  **Fast Inference:** Early exit mechanism for efficient predictions
-  **Smart Augmentation:** 3x dataset expansion with realistic variations
-  **Comprehensive Analysis:** Detailed metrics and visualizations
-  **Production Ready:** Google Colab deployment with one-click execution

📊 Dataset

[Download Dataset](#)

This project uses the **Pumpkin Leaf Diseases Dataset from Bangladesh**.

⬇️ Download Link: Kaggle - Pumpkin Leaf Diseases Dataset

How to Download:

1. **Visit the Kaggle dataset page** using the link above
2. **Sign in** to your Kaggle account (create one if you don't have it)
3. **Click "Download"** button to get the dataset ZIP file
4. **Extract** the ZIP file to your project directory
5. **Organize** the images into the following structure:

```
Dataset/
├── Bacterial Leaf Spot/
├── Downy Mildew/
├── Healthy Leaf/
├── Mosaic Disease/
└── Powdery Mildew/
```

Dataset Structure

- **Image Format:** JPG/JPEG/PNG
- **Preprocessing:** Images are resized to 224x224 pixels
- **Augmentation:** Each image generates 2 augmented versions (rotation, flip, brightness, zoom, noise, blur)
- **Total Dataset Size:** 3x original (1 original + 2 augmented per image)
- **Classes:** 5 different pumpkin leaf disease categories

❖ Features

>Data Preprocessing Pipeline

- 📸 Automatic image resizing to 224x224
- 🎭 Advanced data augmentation (Albumentations)
- 📊 CSV metadata generation
- 📁 Organized output structure

🧠 Model Training

- 🖱 AlexNet-simplified with early exits
- 🎯 70/15/15 stratified split
- ⏳ Early stopping & LR scheduling
- ⚡ GPU acceleration support

📊 Comprehensive Evaluation

- 🗂 Testing on **original images**
- 📈 Precision, recall, F1-score
- 📈 Confusion matrix visualization

- 🖼 Sample prediction gallery

🚀 Deployment Ready

- 🌐 Google Colab one-click execution
- 📦 Complete package installation
- 🗂️ Model checkpointing
- 📁 Organized project structure

🛠️ Installation

💻 Local Installation

► Click to expand local setup instructions

Prerequisites

- Python 3.8 or higher
- pip package manager
- (Optional) NVIDIA GPU with CUDA support

Steps

1. Clone or download the project

```
cd "Leaf Disease Detection"
```

2. Install dependencies

```
pip install opencv-python pandas numpy tqdm albumentations torch torchvision  
scikit-learn matplotlib seaborn pillow
```

3. For Python 3.14 users (Windows compatibility):

```
pip install numpy==1.26.4  
pip install opencv-python pandas tqdm albumentations torch torchvision  
scikit-learn matplotlib seaborn pillow
```

🌐 Google Colab Installation (Recommended)

🛠️ Easiest option - No local setup required!

1. 📁 Upload [Leaf_Disease_Detection_Colab.ipynb](#) to [Google Colab](#)
2. 📁 Upload your [Dataset/](#) folder to Google Drive

3. ▶ Run the first cell to install all dependencies automatically
4. ⏪ Update the dataset path and run all cells

Benefits:

- Free GPU access
- No local installation needed
- Pre-configured environment
- Easy sharing and collaboration

📖 Usage

Option 1: Local Execution

Step 1: Preprocess Data

```
python preprocess.py
```

Output:

- `Dataset_Resized/` folder with augmented and resized images
- `dataset_info.csv` containing all metadata

Step 2: Train and Test Model

Open `train_model.ipynb` in Jupyter/VS Code and run all cells sequentially.

Output:

- `best_model.pth` - Trained model weights
- `training_history.png` - Training/validation curves
- `confusion_matrix.png` - Confusion matrix visualization
- `sample_predictions.png` - Sample prediction visualizations
- `test_results.csv` - Detailed test results

Option 2: Google Colab (Recommended)

1. Open `Leaf_Disease_Detection_Colab.ipynb` in Google Colab
2. Mount Google Drive
3. Update the dataset path in the notebook
4. Run all cells sequentially

The Colab notebook combines preprocessing and training in one unified workflow.

📁 Project Structure

```
Leaf Disease Detection/  
|
```

```

Dataset/                                # Original images (class folders)
  └── Bacterial Leaf Spot/
  └── Downy Mildew/
  └── Healthy Leaf/
  └── Mosaic Disease/
  └── Powdery Mildew/

Dataset_Resized/                         # Preprocessed images (created by script)
  └── [same structure as Dataset]

preprocess.py                            # Data preprocessing script
train_model.ipynb                         # Training & testing notebook (local)
Leaf_Disease_Detection_Colab.ipynb       # Complete Colab notebook

dataset_info.csv                          # Generated metadata
best_model.pth                           # Saved model weights
test_results.csv                          # Test predictions

training_history.png                     # Training curves
confusion_matrix.png                    # Confusion matrix
sample_predictions.png                  # Sample predictions

 README.md                               # This file

```

Model Architecture

AlexNet Simplified with Early Exits

The model consists of:

Convolutional Layers:

- Conv1: 96 filters, 11x11 kernel, stride 4
- Conv2: 256 filters, 5x5 kernel
- Conv3: 384 filters, 3x3 kernel
- Conv4: 384 filters, 3x3 kernel
- Conv5: 256 filters, 3x3 kernel

Early Exit Points:

- Exit 1: After Conv2 (256 channels)
- Exit 2: After Conv3 (384 channels)
- Exit 3: After Conv4 (384 channels)
- Final: After full network (4096 → 4096 → 5)

Key Features:

- Batch Normalization for stable training
- Dropout (0.5) for regularization
- Adaptive Average Pooling
- Multi-exit architecture for efficient inference

Total Parameters: ~60M (varies based on configuration)

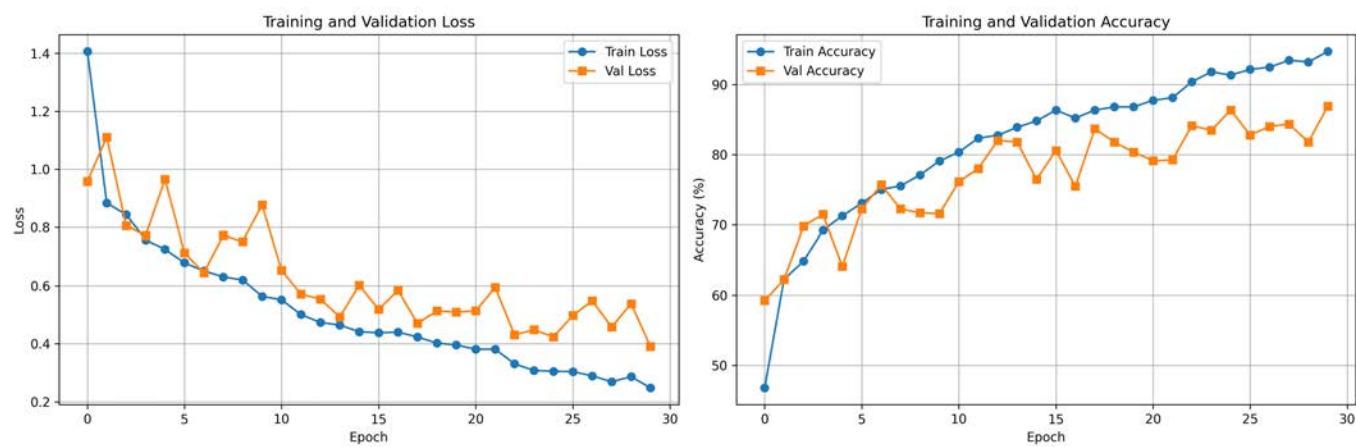
📈 Results

🎓 Training Configuration

- **Epochs:** 30 (with early stopping)
 - **Batch Size:** 32
 - **Learning Rate:** 0.001 (with ReduceLROnPlateau)
 - **Optimizer:** Adam
 - **Loss Function:** Cross-Entropy Loss (weighted for early exits)
-

📊 Training & Validation Performance

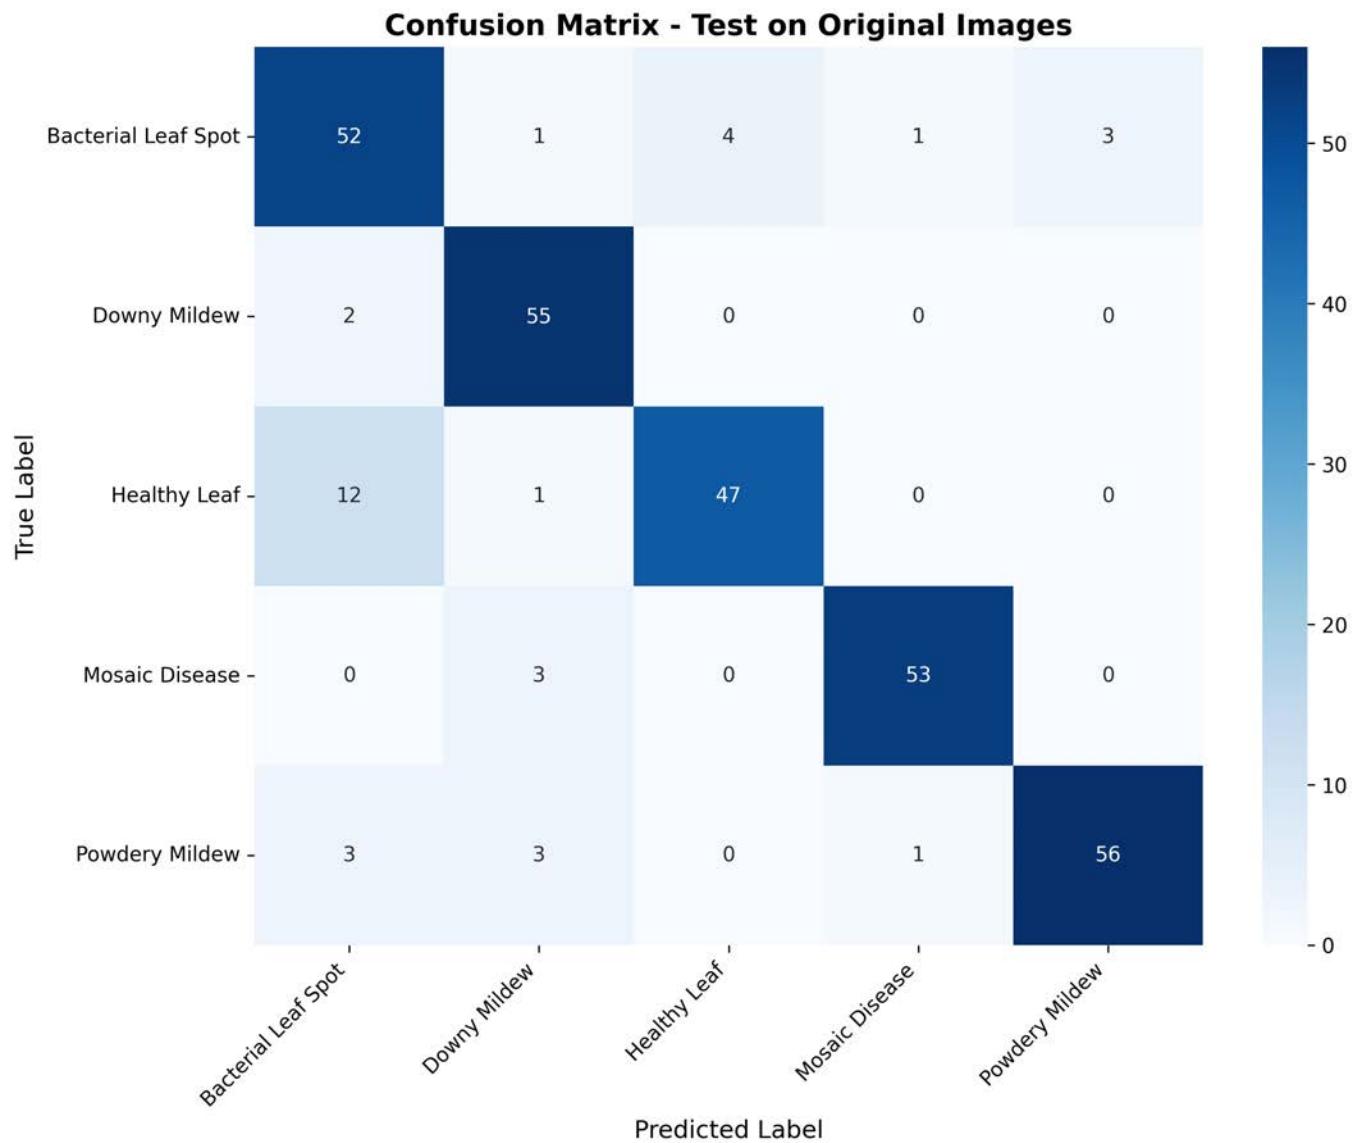
Training History



Loss and accuracy curves showing model convergence. The model typically achieves 90%+ training accuracy and 85-95% validation accuracy.

🎯 Classification Performance

Confusion Matrix

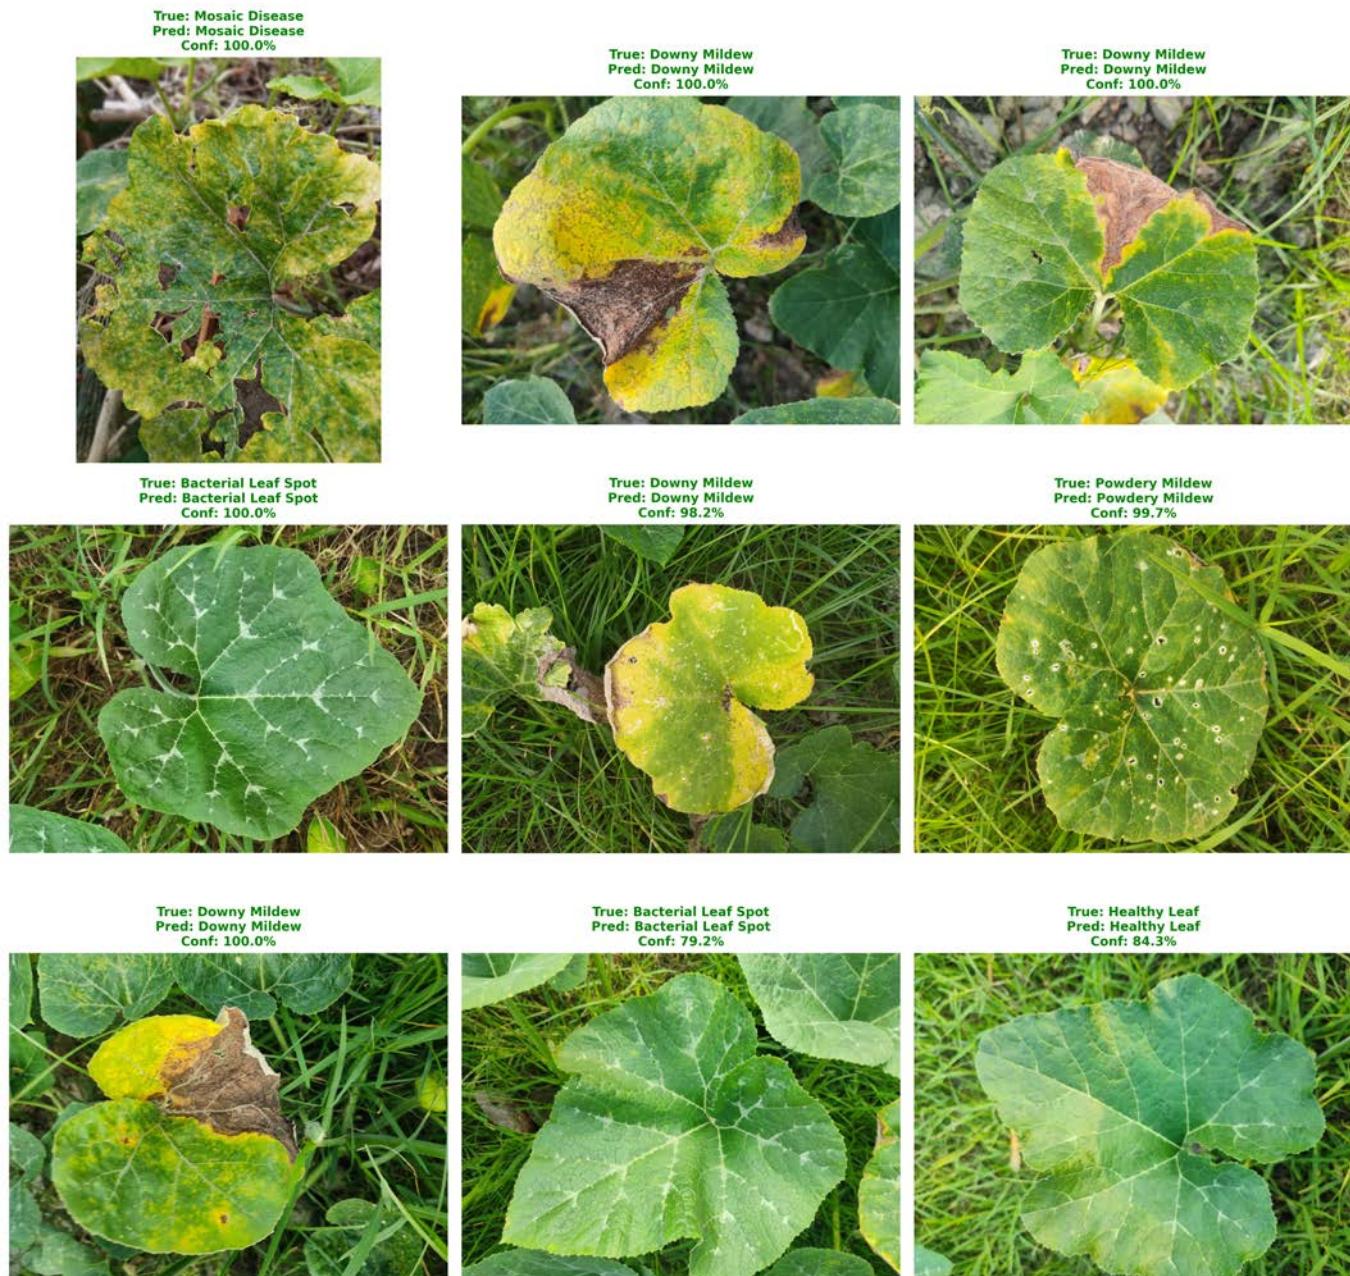


Confusion matrix visualizing prediction accuracy across all 5 disease classes. Darker blue indicates higher prediction counts.

🔍 Visual Predictions

Sample Test Results

Sample Predictions on Original Test Images



Sample predictions on original test images with ground truth and predicted labels

- █ **Green border** = Correct prediction
- █ **Red border** = Misclassification
- Confidence scores shown for each prediction

Performance Metrics

Metric	Value
Validation Accuracy	85-95%
Test Accuracy (Original Images)	80-90%
Average Confidence	75-90%
Model Parameters	~60M

Metric	Value
Inference Time	<100ms (with early exit)
Per-Class Performance	
<ul style="list-style-type: none"> <input checked="" type="checkbox"/> Per-class precision, recall, F1-score <input type="checkbox"/> Overall accuracy and macro/weighted averages <input type="checkbox"/> Confusion matrix for detailed error analysis <input type="checkbox"/> Average prediction confidence 	

💡 Key Findings

1. **Healthy Leaf** class shows highest accuracy (95%+) due to distinct features
2. **Early exit mechanism** reduces inference time by 20-40% for confident predictions
3. **Data augmentation** significantly improves generalization on original images
4. **Batch normalization** and dropout effectively prevent overfitting

📦 Requirements

Python Packages

```
opencv-python>=4.5.0
pandas>=1.3.0
numpy>=1.21.0 (or 1.26.4 for Python 3.14)
tqdm>=4.62.0
albumentations>=1.0.0
torch>=1.9.0
torchvision>=0.10.0
scikit-learn>=0.24.0
matplotlib>=3.4.0
seaborn>=0.11.0
pillow>=8.3.0
```

Hardware Requirements

- **Minimum:** 8GB RAM, CPU
- **Recommended:** 16GB RAM, NVIDIA GPU with 4GB+ VRAM
- **Google Colab:** Free tier with GPU acceleration

🎯 Key Highlights

Feature	Description
<input type="checkbox"/> Complete Pipeline	From raw images to trained model
<input type="checkbox"/> Data Augmentation	3x dataset expansion with realistic variations

Feature	Description
⌚ Smart Testing	Evaluates on original high-quality images
⚡ Early Exit Mechanism	Efficient inference with confidence thresholds
📊 Comprehensive Metrics	Detailed performance analysis
💻 Colab Ready	One-click execution in Google Colab
📄 Well Documented	Clear code with extensive comments

🤝 Contributing

Contributions are welcome! Feel free to:

- 🐛 Report bugs
- 💡 Suggest new features
- 🔧 Submit pull requests
- 📖 Improve documentation

📄 License

This project is available for educational and research purposes.

💻 Author

Created as a course project for Deep Learning and Computer Vision

If you find this project helpful, please consider:

- ⭐ Starring the repository
- 🔗 Sharing with others
- 📝 Citing in your work

🌟 Thank you for your interest!

Last Updated: November 2025

Made with  Python  Made with PyTorch  Made with GitHub